ECM2423 – Artificial Intelligence and Applications Continous Assessment

James Calnan, jdc235@exeter.ac.uk

Deadline: Thursday 16th March

Lecturer: Dr Ayah Helal (a.helal@exeter.ac.uk)

Question 1.1

Describe how you would frame the maze solver as a search problem. In your own words, write a short description of how the maze solver can be seen as a search problem.

In order to frame a maze solver as a search problem you would first have to interpret the maze as a traversable graph, with the start and end points being known. Converting the maze to a traversable graph can be done by iterating over the nodes/vertices and then recording adjacent nodes in order to construct an adjacency list for each node. An adjacency matrix can also be constructed, however in large mazes this would result in a large 2-dimensional array which could put a strain on memory. Once you have constructed an adjacency list for each node, you can then use graph traversal algorithms in order to find a path from the start and end node.

```python
def returnNeighbours(maze, x, y):
    # Initialize an empty list to hold the neighbors of the given
    ↪  cell
    neighbours = []

    # Check the cell above the given cell
    if y - 1 >= 0 and maze[y - 1][x] == "-":
            # If the above cell is within the boundaries of
            ↪  the maze and is a valid path, add it to the
            ↪  list of neighbors
            neighbours.append((x, y - 1))

    # Check the cell below the given cell
    if not (y + 1) == len(maze) and maze[y + 1][x] == "-":
```

```python
                # If the below cell is within the boundaries of
                ↪   the maze and is a valid path, add it to the
                ↪   list of neighbors
                neighbours.append((x, y + 1))

        # Check the cell to the right of the given cell
        if maze[y][x + 2] == "-":
                # If the cell to the right is a valid path, add
                ↪   it to the list of neighbors
                neighbours.append((x + 2, y))

        # Check the cell to the left of the given cell
        if maze[y][x - 2] == "-":
                # If the cell to the left is a valid path, add it
                ↪   to the list of neighbors
                neighbours.append((x - 2, y))

        # Return the list of neighbors
        return neighbours


def buildAdjacencyList(maze):
    bounds = (0, len(maze), 0, len(maze[0]))
    adjacencyList = defaultdict(list)

    for y in range(bounds[1]):
        for x in range(0, bounds[3], 2):
            if maze[y][x] == "-":
                adjacencyList[(x, y)] = returnNeighbours(maze, x,
                ↪   y)

    return adjacencyList
```

Question 1.2

Solve the maze using depth-first search. This question is further divided into parts. This is a guide to help your in the coding process. You do not need to submitted different versions of the code, but just the last updated version.

1. Briefly outline the depth-first algorithm

The depth-first algorithm is a graph traversal algorithm that explores a graph or tree by visiting as far as possible along each branch before backtracking. It can be implemented iteratively using a stack data structure or recursively using a function call stack. The algorithm works by starting at a designated node, visiting its neighbors, and then visiting each neighbor's unvisited neighbors. If a dead end is reached, it backtracks to a previous node with unvisited neighbours (using the stack data structure or function call stack) and continues the search from there until all nodes have been visited or the loop can be terminated once the goal node is reached[3].

While this traversal is happening a solution map can be constructed which will keep track of how each node was reached. This can then be used to backtrack from the goal node to the start node giving you a path.

2. Implement depth-first to solve the maze

The solution of the maze should print the locations in the map that the algorithm will take to reach from the start to the goal. This question can be solved using the file 'maze-Easy.txt'.

See file:

1. maze-Easy-DFS-Solution.txt

Or:

```
# X # # # # # # # # # # # # # # # # # #
# X X X X X X X X X X X X X X X - - # #
# # # # # - # # - # # - # # # X # - - #
# - # # # - # - - # # - # # - X # - # #
# - # - - - # # # # # # # # # X # - - #
# - - - # - - - - # - # # - # X # # # #
# - # # # # - # - - - - - - X X X - #
# - # # - - - # # # # # # # # # X # #
# - - - - # - - # # - - - - - - - X X #
# # # # # # # # # # # # # # # # # # X #
```

```python
def depthFirstSearch(adjacencyList, root, goal):
        # Initialize an empty set to keep track of discovered
        ↪   nodes.
        discovered = set()

        # Initialize a double-ended queue to store nodes to
        ↪   visit.
        S = deque()

        # Add the starting node to the queue.
        S.append(root)

        # Initialize an empty dictionary to keep track of the
        ↪   path from each visited node to its parent.
        cameFrom = {}

        # Initialize a counter to keep track of the number of
        ↪   nodes explored during the traversal.
        nodesExplored = 0

        # While there are nodes in the queue:
        while S:

                # Pop the last node from the queue.
                v = S.pop()

                # Increment the node counter.
                nodesExplored += 1

                # If the current node is the goal, return the
                ↪   dictionary map and the number of nodes
                ↪   explored.
                if v == goal:
                        return cameFrom, nodesExplored

                # If the current node has not been discovered
                ↪   yet:
                if v not in discovered:

                        # Mark it as discovered.
                        discovered.add(v)

                        # For each adjacent node w:
                        for w in adjacencyList[v]:
```

```python
                              # If w has already been
                              ↪  discovered, skip it.
                              if w in discovered:
                                      continue

                              # Otherwise, add w to the queue
                              ↪  and record the path from v to
                              ↪  w.
                              S.append(w)

                              # Set the parent of the neighbour
                              ↪  to the current node
                              cameFrom[w] = v

        # If the goal was not reached, return None for the path
        ↪  and the number of nodes explored.
        return None, nodesExplored
```

Depth-First Search Implementation:[1]

3. Update your algorithm to calculate some statistics about its performance

The number of nodes explored to find the path, the time of the execution, and the number of steps in the resulting path.

Table 1: Statistics for Depth-first Search on maze-Easy.txt (averaged over 5000 runs)

Statistic	Values
Nodes visited	79
Percentage of maze explored	95%
Solution Length	28
Time taken to solve the maze	0.0000358 seconds
Solution percentage	33%

4. Update your algorithm to be generalised to read any maze in the same format. Then calculate the paths and statistics for all 'maze-Medium.txt"/ 'maze-Large.txt" / 'maze-VLarge.txt".

See files:

1. maze-Medium-DFS-Solution.txt

2. maze-Large-DFS-Solution.txt

3. maze-VLarge-DFS-Solution.txt

Table 2: Statistics for Depth-first Search on maze-Medium.txt (averaged over 5000 runs)

Statistic	Values
Nodes visited	2027
Percentage of maze explored	22%
Solution Length	576
Time taken to solve the maze	0.0008746 seconds
Solution percentage	6%

Table 3: Statistics for Depth-first Search on maze-Large.txt (averaged over 1000 runs)

Statistic	Values
Nodes visited	70512
Percentage of maze explored	85%
Solution Length	1051
Time taken to solve the maze	0.0390545 seconds
Solution percentage	1%

Table 4: Statistics for Depth-first Search on maze-VLarge.txt (averaged over 500 runs)

Statistic	Values
Nodes visited	134794
Percentage of maze explored	14%
Solution Length	4050
Time taken to solve the maze	0.0905215 seconds
Solution percentage	0%

Question 1.3

Suggest an improved algorithm for this problem.

1.   Briefly outline another algorithm to solve the maze. That decision should be justified by what you expect to improve in the performance over depth-first search.

One possible algorithm to solve the maze is the A* search algorithm[2]. A* search is an informed greedy search algorithm that uses a heuristic function to guide more efficiently than depth first search. The heuristic function, which in this case will just be the Manhattan distance between two points, estimates the distance between the current state and the goal state, and A* search uses this information to prioritise the exploration of seemingly more promising paths.

By using A* search instead of depth-first search, we can expect to improve the performance of the algorithm in several ways. Firstly, A* search can be more efficient in terms of the number of nodes explored as it doesn't need to explore a path to completion like depth-first search. Secondly, the heuristic can be adjusted to make the algorithm less greedy and more resembling Dijkstra's which will allow for the most optimal solution to be found. Finally, A* search can can be potentially be quicker using a priority queue[4] to dequeue the next node with the lowest heuristic score.

Overall, by using A* search, we can expect to achieve a more optimal solution to the maze compared to depth-first search.

2. Implement the suggested algorithm to solve the maze. It should also calculate the same statistics about its performance as the previous depth-first algorithm. The solution of the maze should print the locations in the map that the algorithm will take to reach from the start to the goal.

See files:

1. maze-Easy-ASTAR-Solution.txt

2. maze-Medium-ASTAR-Solution.txt

3. maze-Large-ASTAR-Solution.txt

4. maze-VLarge-ASTAR-Solution.txt

A* Algorithm

```python
def aStar(adjacencyList, root, goal):
        # Set the heuristic multiplier.
        multiplier = .8

        # Create a dictionary to hold the distances from the root
        ↪  to each node.
        # Initialize the distance to the root to be the heuristic
        ↪  distance.
        distance = defaultdict(lambda: float('inf'))
        distance[root] = heuristic(root, goal, multiplier)

        # Create a dictionary to hold the parent of each node in
        ↪  the shortest path from the root to that node.
        cameFrom = {root: None}

        # Create a binary heap priority queue to store the nodes.
        # Create a priority queue to store the nodes.
        # heap = heapdict()
        prioQueue = PriorityQueue()

        # Enqueue the root with a priority of 0.
        prioQueue.put((0, root))
        # heap[root] = 0

        # Keep track of the number of nodes explored.
        nodesExplored = 0

        # Create a cache for the heuristic values.
        heuristicCache = defaultdict()


        # While there are nodes in the heap.
        while not prioQueue.empty():
                # Extract the node with the lowest priority.
                # current, priority = heap.popitem()
                _, current = prioQueue.get()
                nodesExplored += 1

                # If the current node is the goal, return the
                ↪  shortest path.
                if current == goal:
                        return cameFrom, nodesExplored

                # For each neighbor of the current node.
```

```python
for neighbor in adjacencyList[current]:
        # Calculate the tentative distance from
        ↪   the root to the neighbor through the
        ↪   current node.
        tentative_distance = distance[current] +
        ↪   1 #heuristic(neighbor, current)

        # If the tentative distance is less than
        ↪   the current distance to the neighbor,
        ↪   update the distance.
        if tentative_distance <
        ↪   distance[neighbor]:
                distance[neighbor] =
                ↪   tentative_distance

                # Check if the heuristic value
                ↪   for the neighbor is already
                ↪   cached.
                if neighbor in heuristicCache:
                        # Use the cached value.
                        heuristic_value =
                        ↪   heuristicCache[neighbor]
                else:
                        # Calculate the heuristic
                        ↪   value and cache it.
                        heuristic_value =
                        ↪   heuristic(neighbor,
                        ↪   goal, multiplier)
                        heuristicCache[neighbor]
                        ↪   = heuristic_value

                # Calculate the priority of the
                ↪   neighbor as the sum of the
                ↪   tentative distance and the
                ↪   heuristic distance to the
                ↪   goal.
                priority = tentative_distance +
                ↪   heuristic_value
                # Enqueue the neighbor with the
                ↪   calculated priority.
                # heap[neighbor] = priority
                prioQueue.put((priority,
                ↪   neighbor))
                # Set the parent of the neighbor
                ↪   to the current node.
                cameFrom[neighbor] = current
```

```python
        # If there is no path from the root to the goal, return
        ↪  None for the path and the number of nodes explored.
        return None, nodesExplored


# Heuristic function that calculates the manhatten distance
↪  between two nodes
def heuristic(current, goal, m=1):
        return (abs(current[0] - goal[0]) + abs(current[1] -
        ↪  goal[1])) * m
```

A* Search Algorithm Implementation[2]

4. Discuss based on your results, analysis how the suggested algorithm performed better than the depth-first search algorithm.

Table 5: Statistics for A* (.8 Multiplier) on maze-Easy.txt

Statistic	Values
Nodes visited	35
Percentage of maze explored	42%
Solution Length	28
Time taken to solve the maze	0.0000524 seconds
Solution percentage	33%

Table 6: Statistics for A* (.8 Multiplier) on maze-Medium.txt (averaged over 5000 runs)

Statistic	Values
Nodes visited	456
Percentage of maze explored	5%
Solution Length	322
Time taken to solve the maze	0.0007493 seconds
Solution percentage	3%

Table 7: Statistics for A* (.8 Multiplier) on maze-Large.txt (averaged over 1000 runs)

Statistic	Values
Nodes visited	41752
Percentage of maze explored	50%
Solution Length	975
Time taken to solve the maze	0.069606 seconds
Solution percentage	1%

Table 8: Statistics for A* Search (.8 Multiplier) on maze-VLarge.txt (averaged over 500 runs)

Statistic	Values
Nodes visited	74999
Percentage of maze explored	8%
Solution Length	3692
Time taken to solve the maze	0.1380474 seconds
Solution percentage	0%

Results were obtained running DFS and A* on a Ryzen 5 3600.

Based on the results obtained it can be seen that the A* algorithm is required to visited much fewer nodes than Depth-First Search in order to reach the goal node. However, despite needing to visited fewer nodes due to the overhead and more complex operations in the A* algorithm (ordering priority queue, storing heuristic values) the execution time is always longer. This can also be due to the nature of how I coded the adjacency list. When I create the adjacency list the neighbouring nodes are placed into the dictionary in this order: above the current node, below the current node, to the right of the current node and to the left of the current node. This means that when the neighbouring nodes are added to the stack in the depth first search it is the node below takes precedence over the nodes to the right and the left. This works well for the inputted mazes as the search starts from the top of the maze and works its way to the bottom. When this ordering is changed to below, right, left, above the performance of the algorithm suffers:

Table 9: DFS on maze-VLarge.txt with different adjacency list ordering

Statistic	Values
Nodes visited	778817
Percentage of maze explored	84%
Solution Length	6132
Time taken to solve the maze	0.60223 seconds
Solution percentage	0%

In the above example the depth first search visited 84% of the maze instead of just 14% with the optimal neighbour ordering. The solution length is also much longer at 6132 instead of 4050 with the original neighbour ordering.

Table 10: Comparison of solution lengths

Maze File	Algorithm	Solution length
maze-Easy.txt	A*	28
	DFS	28
maze-Medium.txt	A*	322
	DFS	576
maze-Large.txt	A*	975
	DFS	1051
maze-VLarge.txt	A*	3692
	DFS	4050

As can be seen from the table above (Table 10), the A* algorithm was able to find a solution which has a shorter length on all but the easy maze (which both algorithms found a value of 28). Based on these results the A* algorithm can be seen to performed better than the depth first search algorithm. Additionally, the A* algorithm wont suffer from the same ordering of the adjacency list problem that Depth-First Search suffers from due to the use of a priority queue.

Table 11: Comparison of maze exploration

Maze File	Algorithm	Nodes Visited	% of maze visited
maze-Easy.txt	A*	35	42%
	DFS	79	95%
maze-Medium.txt	A*	456	5%
	DFS	2027	22%
maze-Large.txt	A*	41,752	50%
	DFS	70,512	85%
maze-VLarge.txt	A*	74,999	8%
	DFS	134,794	14%

From the statistics in the table above (Table 11) it can be observed that the A* algorithm was required to visit less nodes than Depth-First Search in order to find a solution.

References

[1] James Calnan. Maze-generation-python. https://github.com/jamescalnan/Maze-Generation-Python, 2023. GitHub repository.

[2] Wikipedia. A* search algorithm, 2023.

[3] Wikipedia. Depth-first search, 2023.

[4] Wikipedia. Priority queue, 2023.